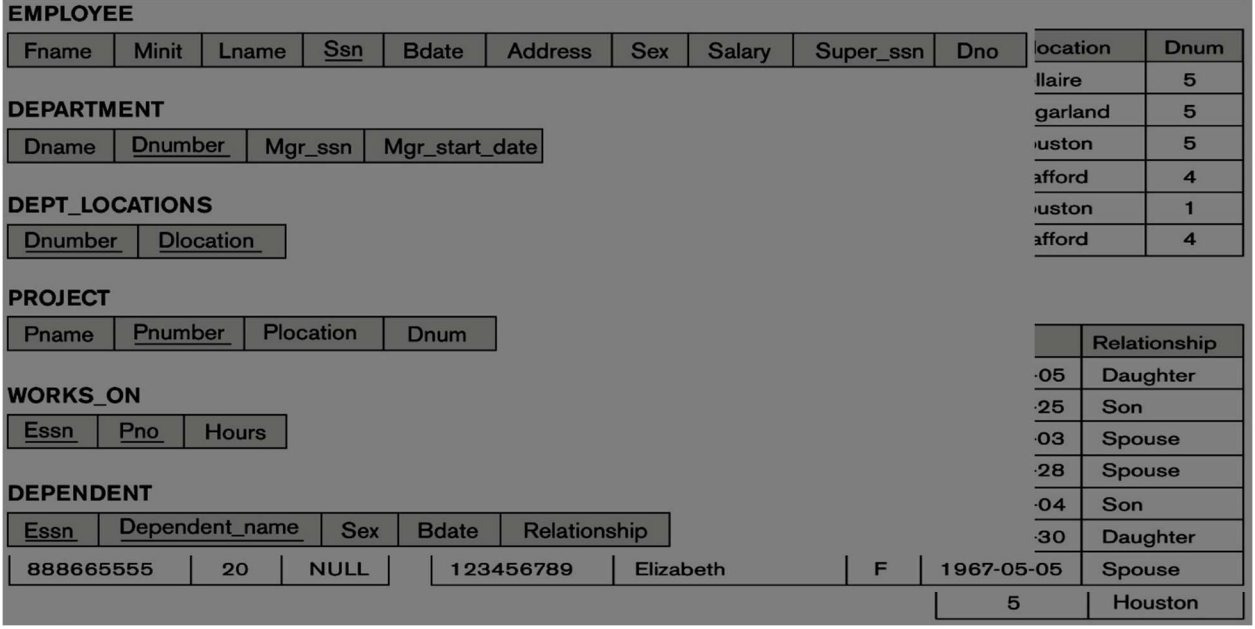

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date

DEPT_LOCATIONS

Dnumber	Dlocation

PROJECT

Pname	Pnumber	Plocation	Dnum

WORKS_ON

Essn	Pno	Hours

DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
888665555	20	NULL	123456789	Elizabeth	F	1967-05-05	Spouse

location	Dnum
llaire	5
garland	5
uston	5
afford	4
uston	1
afford	4

	Relationship
05	Daughter
25	Son
03	Spouse
28	Spouse
04	Son
30	Daughter
5	Houston

School: Computer Science
Institution: University of Windsor
Term: Fall 2018
Course: 03-60-315-1 : Database Management Systems
Instructor: Dr. C. I. Ezeife
Assignment #3: Total: 50 marks
Handed Out:Thurs. Oct. 25, 2018; Due Thurs Nov. 15, 2018

Objective of Assignment: To test on knowledge and use of relational database query languages SQL and relational algebra for implementing relational databases.
Scope: Assignment covers materials from Chapters 6, 7 and 8 of book discussed in class.
Electronic Assignment Submission: Done through http://blackboard.uwindsor.ca
Marking Sheme : The mark for each of the questions is indicated beside each question.
Academic Integrity Statement: Remember to submit only work that is yours and include the following confidentiality agreement and statement at the beginning of your assignment.

CONFIDENTIALITY AGREEMENT & STATEMENT OF HONESTY

I confirm that I will keep the content of this assignment/examination confidential.
I confirm that I have not received any unauthorized assistance in preparing for or doing this assignment/examination. I confirm knowing that a mark of 0 may be assigned for copied work.

_____Lyndon Renaud_____ _____Lyndon Renaud_____
Student Signature Student Name (please print)
____104 566 776_____ _____11/14/2018_____
Student I.D. Number Date

Marking Scheme : The mark for each question and sub question is shown with the question below. Place your solutions in tables where possible.

For office Use only

Question	Mark
1	/15

2	**/10**
3	**/5**
4	**/5**
5	**/5**
6	**/10**
Total	**/50**

CHAPTER 6: Basic SQL

1. Given a database state of the AIRLINE RESERVATION database shown in Figure 1.1, with schema shown in Figure 1.2.

(Total for que 1 is 15 marks)

Fig 1.1: An Example Database State of Parts of the Airline Database of Fig 5.8 of book

AIRPORT

Airport_Code	Name	City	State
YYZ	Pearson Airport	Toronto	ON
YTZ	Island Airport	Toronto	ON
YQT	Thunder Bay Airport	Thunder Bay	ON
YQG	Windsor Airport	Windsor	ON
YYC	Calgary Airport	Calgary	AB

FLIGHT

Flight_number	Airline	Weekdays
WJ250	WestJet	3
WJ261	WestJet	4
AC275	Air Canada	2
AC300	Air Canada	3
AC320	Air Canada	2
PA233	Porter Airline	5
PA280	Porter Airline	4

FLIGHT_LEG

Flight_number	Leg_number	Departure_airport_code	Scheduled_departure_time	Arrival_airport_code	Scheduled_arrival_time
WJ250	1	YQG	08.00	YYZ	09.00
WJ250	2	YYZ	12.00	YYC	17.00
AC275	1	YQT	10.00	YYZ	11.00

AC275	2	YYZ	13.00	YQG	14.00
PA233	1	YQG	16.00	YYZ	17.00
PA280	1	YYZ	11.00	YQG	12.00

LEG_INSTANCE

Flight_number	Leg_number	FDate	Num_available_seats	Airplane_id	Departure_airport_code	Departure_time	Arrival_airport_code	Arrival_time
WJ250	1	15-JAN-17	200	Boeing_777	YQG	8.00	YYZ	8.50
WJ250	2	15-JAN-17	200	Boeing_777	YYZ	12.30	YYC	16.50
AC275	1	10-JAN-17	60	Airbus_200	YQT	10.00	YYZ	11.00
AC275	2	10-JAN-17	60	Airbus_200	YYZ	13.00	YQG	14.00
PA233	1	30-JAN-17	120	Boeing_737	YQG	16.00	YYZ	17.00
PA280	1	20-JAN-17	120	Boeing_737	YYZ	11.00	YQG	12.00

FARE

Flight_number	Fare_code	Amount	Restrictions
WJ250	F1	$250	None
WJ250	F2	$350	None
AC275	F1	$275	Non-refundable
AC275	F2	$300	Non-refundable
PA233	F3	$150	None
PA280	F3	$150	None

SEAT_RESERVATION

Flight_number	Leg_number	FDate	Seat_number	Customer_name	Customer_phone
WJ250	1	15-JAN-17	20A	Mariane Mooer	519-253-3000
WJ250	2	15-JAN-17	13D	Mariane Mooer	519-253-3000
WJ250	1	15-JAN-17	6A	Tony Appa	226-253-4000
WJ250	2	15-JAN-17	7D	Tony Appa	226-253-4000
WJ250	1	15-JAN-17	6B	Karen Appa	226-253-4000
WJ250	2	15-JAN-17	7C	Karen Appa	226-253-4000
AC275	1	10-JAN-17	5C	Mark Black	226-212-5341

AC275	2	10-JAN-17	14B	Mark Black	226-212-5341
PA233	1	30-JAN-17	5A	Peter Opo	519-444-1234
PA280	1	20-JAN-17	16D	Chris Natta	226-123-5555

Fig 1.2: Schema of the Airline Reservation database of Figure 1.1

AIRPORT

Airport_Code	Name	City	State

FLIGHT

Flight_number	Airline	Weekdays

FLIGHT_LEG

Flight_number	Leg_number	Departure_airport_code	Scheduled_departure_time	Arrival_airport_code	Scheduled_arrival_time

LEG_INSTANCE

Flight_number	Leg_number	FDate	Num_available_seats	Airplane_id	Departure_airport_code	Departure_time	Arrival_airport_code	Arrival_time

FARE

Flight_number	Fare_code	Amount	Restrictions

SEAT_RESERVATION

Flight_number	Leg_number	FDate	Seat_number	Customer_name	Customer_phone

i. List all the referential integrity constraints that should hold on the database schema? (2.5 marks)

ii. Write appropriate SQL DDL statements to define the database with the integrity constraints and store in a text file called userid_airlineschema.sql. Attach this file or also show it in your script file of (v) using more file.sql command before or after running sqlplus. Do the same for the files in (iii) and (iv). (2.5 marks)

iii. To insert the data in the database tables, also write appropriate SQL DML instructions in a text file called userid_airlinedata.sql. (2.5 marks)

iv. To remove any inserted data and destroy all created tables in the airline reservation database, write appropriate SQL DML and DDL statements in a text file called userid_airlinedroptable.sql to first delete all data in the tables and then drop the tables. (2.5 marks)

v. Using Oracle Sqlplus, implement this database design by creating all the tables with the integrity constraints using the SQL DDL you defined in (ii) above. You can create all these SQLfor the 6 tables by running your .sql file at the SQL prompt with the command: @userid_airlineschema.sql. After creating your tables successfully, you load your data with the .sql file you created in (iii) above by running @userid_airlinedata.sql. If there are errors and you need to correct them, you might want to delete the tuples and drop the tables first using the .sql file you created in (iv) above as with @userid_airlinedroptable.sql before re-creating the schema and re-loading the data.Then, using a script file, show the contents of all 6 tables in the database by selecting * from each of the tables and saving on script file called username_assn3que1.txt. You can do this using the following sequence of Unix/Linux commands after you have created the database and inserted data. (5 marks)

(Note: remember to create the entity tables with primary keys before the relationship tables that reference them through foreign key attributes. When inserting data, do the same. If you need to delete the data and tables at any time, go in the reverse order (that is, delete the tuples that reference a primary key attribute tuple in another table, before deleting the parent primary keyed tuple))

```
>script   username_assn3que1.txt
>sqlplus <username>
>password
sqlplus>  select *  from AIRPORT;   //repeat this instruction for each table
sqlplus> exit                        //to exit sqlplus
exit                                 // to exit and create script file
```

**Now attach the saved log of your session that is in username_assn3que1.txt with an inclusion in this script file of all the 3 .sql files in questions (ii), (iii) an (iv) or the attachment of those files as your solution.

Solution 1 (i) (mark: 2.5)

Flight_number of FLIGHT_LEG references Flight_number of FLIGHT. (Flight_number, Leg_number) of LEG_INSTANCE references (Flight_number, Leg_number) of FLIGHT_LEG Flight_number of FARE references Flight_number of FLIGHT. (Flight_number, Leg_number, FDATE) of SEAT_RESERVATION references (Flight_number, Leg_number, FDATE) of LEG_INSTANCE.

Solution 1 (ii): (mark: 2.5)

Attached in file renau11s_airlineschema.sql

Solution 1 (iii): (mark: 2.5)

Attached in file renau11s_airlinedata.sql

Solution 1 (iv): (mark: 2.5)

Attached in file renau11s_airlinedroptable.sql

1 (v). (5 marks) for the script file showing correct interaction with Oracle Sqlplus creating and loading data in these 6 tables.

2. Specify the following 5 queries in SQL on the airline database schema of Figure 1.1.
(Total for que 2 is 10 marks)

i. List all your 5 SQL queries in the table below first in SQL. (5 marks)
ii. Implement the answering of your 5 SQL queries in 2(i) using Sqlplus and the same database you created in question 1, providing your execution and answers to this questions in a script file called username_assn3que2.txt. (5 marks)

(a) Retrieve the names and flight numbers of all customers who have taken a flight departing from Windsor.
(b) Retrieve the flight number and flight legs of all flights that have more than one leg.
(c) For each flight, retrieve the flight number, airline, flight date, and number of customers in the flight.
(d) Retrieve the name, seat number, and flight number of each customer in all Air Canada flights. (e) Retrieve the names and flight numbers, departure and arrival airports of all customers who did not fly in a flight that is more than the first leg.

Solution 2 (i): SQL Queries(5 marks) and 2(ii) Results (5 marks)

a)

```
SELECT sr.customer_name, sr.flight_number
FROM seat_reservation sr, airport, leg_instance li
WHERE airport.city='Windsor' and li.departure_airport_code=airport.air-
port_code and sr.flight_number=li.flight_number and sr.leg_num-
ber=li.leg_number;
```

b)

```
SELECT flight_number, COUNT(leg_number) flight_legs
FROM flight_leg
HAVING COUNT(leg_number) > 1
GROUP BY flight_number
```

c)

```
SELECT sr.flight_number, airline, fdate, count(customer_name) customers
FROM seat_reservation sr, flight
WHERE flight.flight_number=sr.flight_number
GROUP BY sr.flight_number, airline, fdate;
```

d)

```
SELECT sr.customer_name, sr.seat_number, sr.flight_number
FROM seat_reservation sr, flight
WHERE airline='Air Canada' and sr.flight_number=flight.flight_number;
```

e)

```
SELECT distinct sr.customer_name, sr.flight_number, departure_air-
port_code, arrival_airport_code
FROM seat_reservation sr, flight_leg fl
WHERE fl.flight_number=sr.flight_number and fl.leg_number=sr.leg_number
and sr.leg_number=1
```

2 (ii). (5 marks) distributed as: 2.5 marks for the script file showing correct interaction with Oracle Sqlplus posing these queries; and 2.5 marks for the correctly posing the queries and retrieving correct results.

3. Write four SQL update statements to do the following updates on the database schema shown in Figure 1.2. Show the affected tables after update through script file in sqlplus and in a script file created as before and named username_assn3que3.txt. (5 marks)

(Total for que 3 is 5 marks)

(a) Insert a new airport <'YTT', 'Tintin Airport', 'Tintin', 'ON'> in the database.

(b) Change the Customer_phone of customer 'Tony Appa' to 519-253-4001.

(c) Insert a new flight <'AC331','Air Canada', 2>.

(d) Delete all reservation records for the customer whose name is 'Mariane Mooer '.

Solution 3 (i): (5 marks)

a)

```
INSERT INTO airport
VALUES('YTT', 'Tintin Airport', 'Tintin', 'ON');
```

b)

```
UPDATE seat_reservation
SET customer_phone='519-253-4001'
WHERE customer_name='Tony Appa';
```

c)

```
INSERT INTO flight
VALUES('AC331', 'Air Canada', 2);
```

d)

```
DELETE from seat_reservation
WHERE customer_name='Mariane Mooer';
```

CHAPTER 7: More SQL: Complex Queries, Triggers, Views, and Schema Modification

4. (i) Write the following 2 queries in SQL on the database schema of Figure 1.2 using EXISTS or NOT EXISTS as appropriate.

(2.5 marks)

(ii) Implement the answering of your 2 queries in 4(i) using Sqlplus and the same database you created in question 1, providing your execution and answers to this question in a script file called username_assn3que4. (2.5 marks)

(Total for que 4 is 5 marks)

(a) Retrieve the customer names and flight number of all 2-leg flyers (customers who have flights with 2 legs in all their flights).
(b) Retrieve the customer names and flight number of all customers who do not have any 2-leg flight in any of their flights.

Solution 4 (i): (2.5 marks)

a)
```
SELECT s1.customer_name, s1.flight_number
FROM seat_reservation s1
WHERE NOT EXISTS
(SELECT s2.flight_number
FROM seat_reservation s2
WHERE s2.customer_name=s1.customer_name and leg_number=1);
```

b)
```
SELECT s1.customer_name, s1.flight_number
FROM seat_reservation s1
WHERE NOT EXISTS
(SELECT s2.flight_number
FROM seat_reservation s2
WHERE s2.customer_name=s1.customer_name and leg_number=2);
```

4 (ii). (2.5 marks) distributed as: 0.5 marks for the script file showing correct interaction with Oracle Sqlplus posing these queries; and 2 marks for the correctly posing the queries and retrieving correct results.

5. In SQL, specify the following 3 queries on the COMPANY database of Figure 5.5 using the concept of nested queries and the concepts described in chapter 7. (Total for que 5 is 5 marks)

a. Retrieve the names of all employees who work in the department that has the employee with the lowest salary among all employees.
b. Retrieve the names of all employees whose supervisor's supervisor has '333445555' for Ssn.
c. Retrieve the names of employees who make at least $10,000 more than the employee who is paid the least in the company.

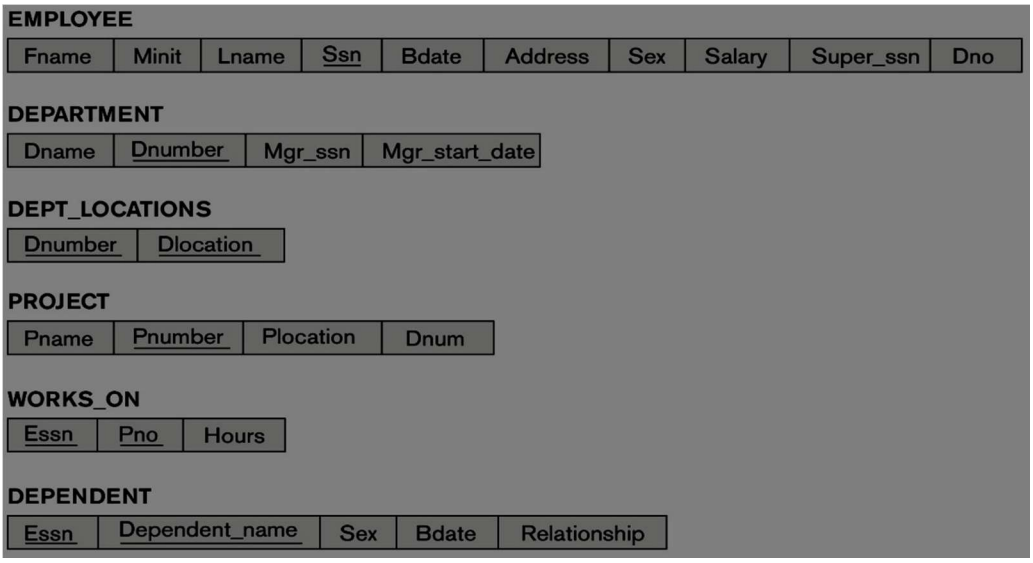

Solution 5: **(5 marks)**

a)

```
SELECT Fname, Lname
FROM employee, department
WHERE dno=dnumber AND dnumber IN
(SELECT dnumber
FROM department, employee
WHERE dnumber=dno AND salary IN
(SELECT salary FROM employee
WHERE salary in (select min(salary) from employee)));
```

b)

```
SELECT Fname, Lname FROM employee
WHERE super_ssn in
(SELECT ssn FROM employee
WHERE super_ssn=333445555);
```

c)

```
SELECT Fname, Lname FROM employee
WHERE salary > (select min(salary) from employee)+10000;
```

CHAPTER 8: THE RELATIONAL ALGEBRA AND RELATIONAL CALCULUS

6. Specify the following 5 queries on the COMPANY relational database schema shown in Figure 5.5, using the relational operators discussed in chapter 8. Also show the result of each query as it would apply to the database state of Figure 5.6. (Total for que 6 is 10 marks)

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

(i) List the names of employees who have a dependent with the same first name as themselves.

(ii) Find the names of employees that are directly supervised by 'James Borg'.

(iii) For each project, list the project name and the total hours per week (by all employees) spent on that project.

(iv) Retrieve the names of employees who work on every project.

(v) Retrieve the maximum salary of all female employees.

Solution 6: (10 marks)

Result of these queries are stored in table RESULT

i)

$$RESULT \leftarrow \pi_{Fname,\ Lname}(\sigma_{Fname=Dependent\_name}(EMPLOYEE \bowtie_{Essn=Ssn} DEPENDENT))$$

Fname	Lname

No employees have a dependent with the same first name as themselves.

ii)

$$jamesSSN(Super\_ssn) \leftarrow \pi_{Ssn}(\sigma_{Fname='James'\ and\ Lname='Borg'}(EMPLOYEE))$$
$$RESULT \leftarrow \pi_{Fname,\ Lname}(EMPLOYEE \div jamesSSN)$$

Fname	Lname
'Franklin'	'Wong'
'Jennifer'	'Wallace'

Only two employees are directly supervised by James Borg

iii)

RESULT <- $\pi_{pname,\ total\_hours}(\rho_{R(pno,\ pname,\ total\_hours)}(_{Pno,\ pname}\Im_{SUM\ Hours}(WORKS\_ON \bowtie_{pno=pnumber} PROJECT)))$

pname	total_hours
'ProductX'	52.5
'ProductY'	75.0
'ProductZ'	50.0
'Computerization'	55.0
'Reorganization'	25.0
'Newbenefits'	55.0

iv)

emp_works_on <- $(WORKS\_ON \bowtie_{pno=pnumber} PROJECT)$

proj_number(pno) <- $\pi_{pnumber}$ (PROJECT)

RESULT <- $\pi_{Fname,\ Lname}$ (emp_works_on ÷ proj_number)

Fname	Lname
Nobody works on all projects	

v)

salary_by_sex <- $\rho_{R(sex,\ max\_salary)}\ (_{sex}\Im_{MAX\ Salary})$

Result <- $\sigma_{sex='F'}$ (salary_by_sex)

sex	max_salary
'F'	43000